



互联网应用下的大规模在线学习算法

陈景东 2013-7-25

大纲

- 通用在线学习算法
 - 在线学习算法的应用背景
 - SGD、mini-batch算法简介 (Andrew NG)
 - SGD算法的收敛特性 (T. Zhang)
- SGD稀疏化方法
 - 稀疏化方法概要
 - Truncate Gradient
 - Regularized Dual Averaging
 - TFRL-Proximal
- 大规模在线学习系统实践
 - 节省内存
 - 有效地模型评估
 - 其它问题的探讨

在线学习算法应用示例

□ 视频分析



□ 场景复杂、多变，如何实现运动目标检测

□ CTR预估

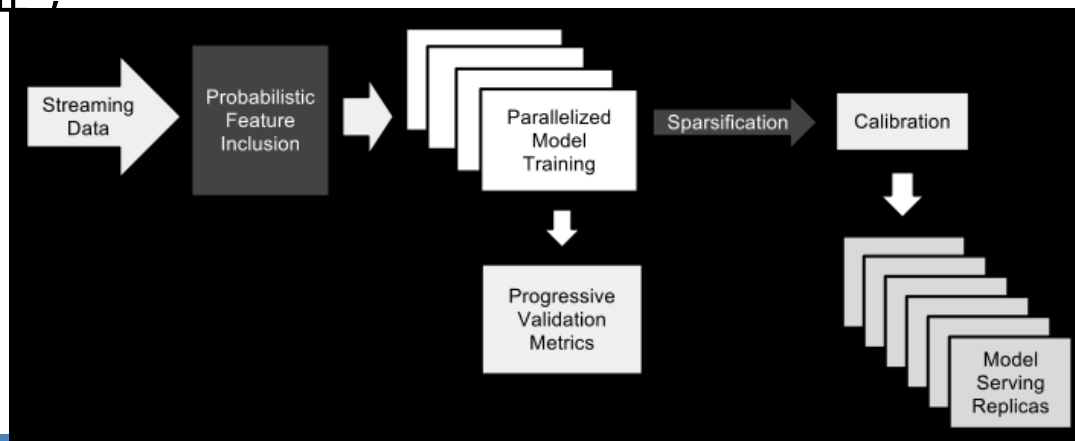
□ 广告库中选出用户最感兴趣的样本；

□ 数以万亿计的样本及特征；

□ 算法适用环境

□ 环境复杂多变；

□ 持续动态输入的大量样本；



Batch gradient descent

□ 以线性回归为例

$$h_{\theta}(x) = \sum_{j=0}^n \theta_j x_j$$

$$J_{train}(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$


Repeat {

$$\theta_j := \theta_j - \alpha \left(\frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} \right)$$

(for every $j = 0, \dots, n$)

}

$m = 300,000,000$



□ 每迭代一次需要读取所有样本，计算量大


Stochastic gradient descent

□ 以线性回归为例

$$\text{cost}(\theta, (x^{(i)}, y^{(i)})) = \frac{1}{2} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

$$J_{\text{train}}(\theta) = \frac{1}{2m} \sum_{i=1}^m \text{cost}(\theta, (x^{(i)}, y^{(i)}))$$

1. Randomly shuffle dataset.
2. Repeat {
 for $i := 1, \dots, m$ {
 $\theta_j := \theta_j - \alpha (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$
 (for $j = 0, \dots, n$)
 }
}



- ## □ 学习率 α 通常为一个充分小的常数， α 也可以按照时间逐步衰减

$$\text{(E.g. } \alpha = \frac{\text{const1}}{\text{iterationNumber} + \text{const2}} \text{)}$$

- ## □ 若样本不完全随机？

Mini-batch gradient descent

Say $b = 10$, $m = 1000$.

Repeat {

➤ for $i = 1, 11, 21, 31, \dots, 991$ {

$$\Rightarrow \theta_j := \theta_j - \alpha \frac{1}{10} \sum_{k=i}^{i+9} (h_{\theta}(x^{(k)}) - y^{(k)}) x_j^{(k)}$$

(for every $j = 0, \dots, n$)

}

}

$\mu = 300,000,000$

□ 对比SGD

□ 梯度方向更稳定

□ 收敛速度更慢

Standard SGD

□ 目标函数

$$Q_\lambda(w) = E_{X,Y} \phi(w^T X, Y) + \frac{\lambda}{2} \|w\|_2^2,$$

□ Standard SGD

Initialize \hat{w}_0

for $t = 1, 2, \dots$

 Draw (X_t, Y_t) randomly from D .

 Update \hat{w}_{t-1} as:

$$\hat{w}_t = \hat{w}_{t-1} - \eta_t S^{-1} (\lambda \hat{w}_{t-1} + \phi'_1(w_{t-1}^T X_t, Y_t) X_t)$$

end

□ 其中 $\phi'_1(p, y) = \frac{\partial}{\partial p} \phi(p, y)$

Standard SGD

- 事先将学习率固定为一个比较小的常数，是一个简单实用的办法；
- 在数据集上多几次迭代，SGD算法更接近batch最优解；
- 数据集上的迭代次数与正则化效应密切相关
 - 迭代次数越少正则化效应越强
 - Go through the data oncethe variance of the estimator is relatively small
 - Go through the data multiple time.... The variance increase;
 - SGD算法在数据集上1次（少量的几次）即可收敛,有比较好的正则化效应，这是SGD算法高效的重要原因

Averaged SGD

□ Averaged SGD

Initialize \hat{w}_0

Let $\hat{v}_0 = 0$ and $r_0 = 0$

for $t = 1, 2, \dots$

 Draw (X_t, Y_t) randomly from D .

 Update \hat{w}_{t-1} as:

$$\hat{w}_t = \hat{w}_{t-1} - \eta_t S^{-1} (\lambda \hat{w}_{t-1} + \phi'_1(w_{t-1}^T X_t, Y_t) X_t)$$

$$r_t = r_{t-1} + \eta_t - \frac{AM^2\eta_t^2}{2}$$

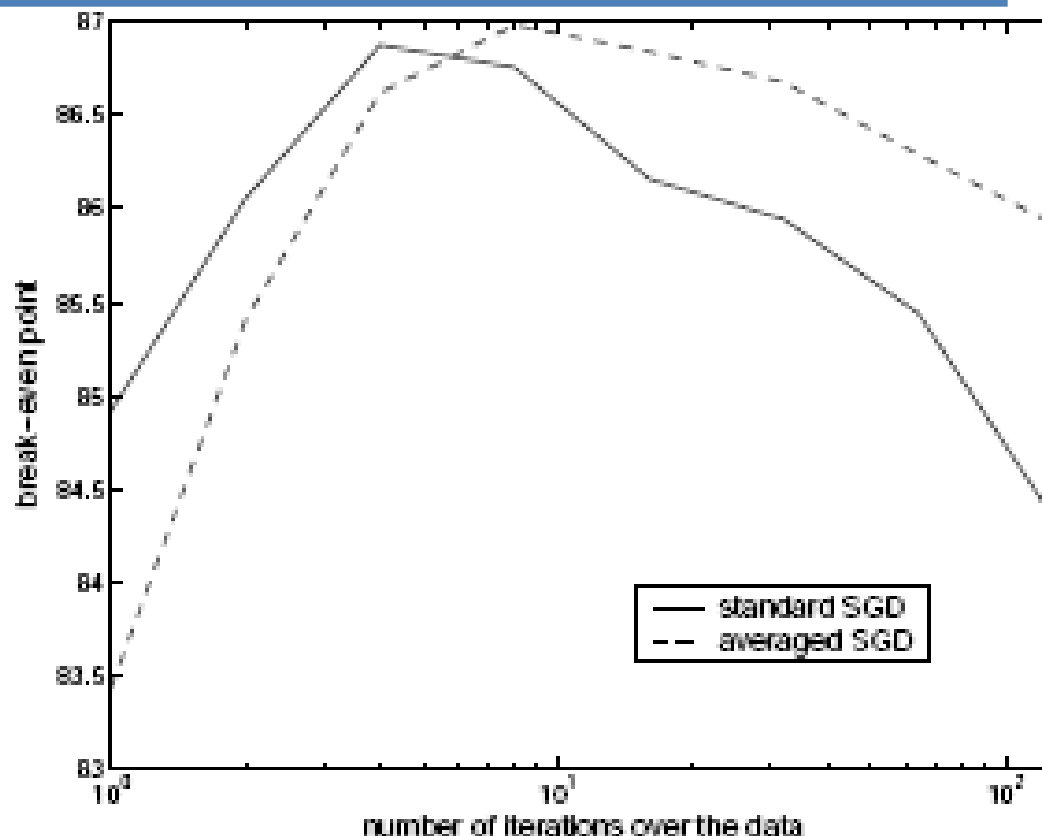
$$\hat{v}_t = \frac{r_{t-1}}{r_t} \hat{v}_{t-1} + \frac{r_t - r_{t-1}}{r_t} \hat{w}_{t-1}$$

end

□ Averaged predictor is more stable than non-averaged

算法对比

- SGD算法收敛速度快（约10步）；
- Averaging并没有带来收益，SGD源于学习率比较小；
- Early stopping与正则化效果相当；



收敛性证明思路

- In separable case , stochastic update with quadratically smoothed SVM behaves similar to the perceptron algorithm
- Using a technique in the online learning literature to extend the traditional proof of perceptron mistake bounds




Figure 2. Performance of perceptron and SGD with SVM loss versus the number of iterations over the data. non-averaged methods: 'solid lines'; averaged methods: 'dashed lines'; perceptron: 'x', SGD : 'o'

大纲

- 通用在线学习算法
 - 在线学习算法的应用背景
 - SGD、mini-batch算法简介 (Andrew NG)
 - SGD算法的收敛特性 (T. zhang)
- SGD稀疏化方法
 - 稀疏化的必要性
 - Truncate Gradient
 - Regularized Dual Averaging
 - TFRL-Proximal
- 大规模在线学习系统实践
 - 节省内存
 - 有效地模型评估
 - 其它问题的探讨

参数稀疏化方法概述

- 参数稀疏化的必要性
 - 内存的限制（训练过程中）；
 - 快速预估服务的需求（训练结束）；
- 一些无效的正则化方法
 - L1 正则化；
 - 简单的四舍五入；
- 有效地稀疏化方法
 - Truncate Gradient (J. Langford et al, JMLR 2009)
 - Regularized Dual Averaging(L. Xiao NIPS 2009)
 - TFRL-Proximal(H. B. McMahan et al, ICML 2013)

Truncate Gradient

- Simple Coefficient Rounding

$$f(w_i) = T_0(w_i - \eta \nabla_1 L(w_i, z_i), \theta),$$

- where $T_0(v_j, \theta) = v_j I(|v_j| < \theta)$

- Too aggressive , a large number of infinitesimal steps;

- Truncate Gradient

$$f(w_i) = T_1(w_i - \eta \nabla_1 L(w_i, z_i), \eta g_i, \theta),$$



- Where

$$T_1(v_j, \alpha, \theta) = \begin{cases} \max(0, v_j - \alpha) & \text{if } v_j \in [0, \theta] \\ \min(0, v_j + \alpha) & \text{if } v_j \in [-\theta, 0] \\ v_j & \text{otherwise} \end{cases}.$$

Empirical Results

- Left: the amount of features left after sparsification for each dataset without 1% performance loss
- AUC tolerated to 1.3%, achieve 91.4%

- the ratio of AUC with and without sparsification.
- AUC improve due to removal of some irrelevant features



Empirical Results

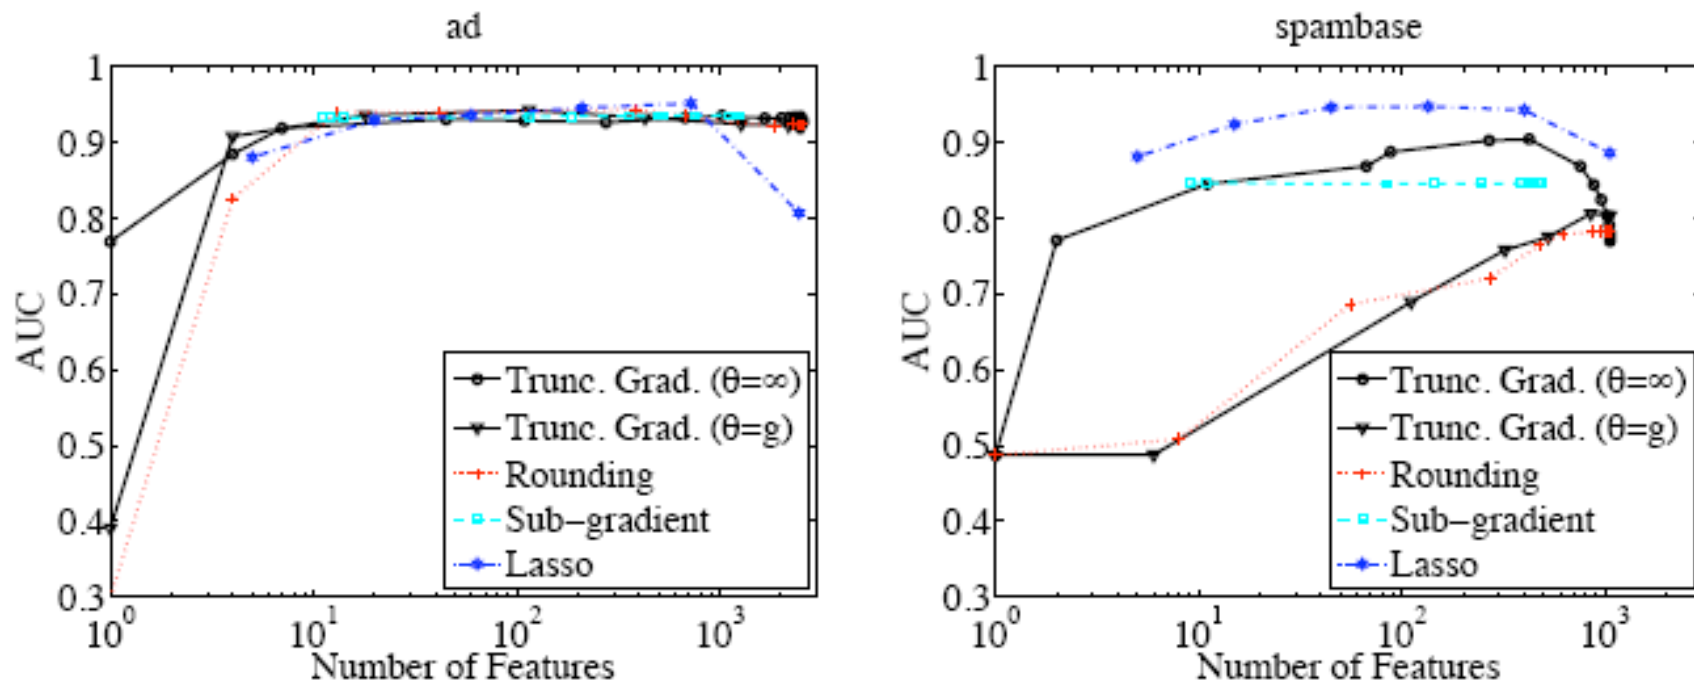


Figure 2: Comparison of the five algorithms in two sample UCI datasets.

- Truncate Gradient, efficient sparsification technique with strong theoretical guarantees

Regularized Dual Averaging

Algorithm 1 Regularized dual averaging (RDA) method

input:

- a strongly convex function $h(w)$ with modulus 1 on $\text{dom}\Psi$, and $w_0 \in \mathbf{R}^n$, such that

$$w_0 = \arg \min_w h(w) \in \text{Arg} \min_w \Psi(w). \quad (4)$$

- a pre-determined nonnegative and nondecreasing sequence β_t for $t \geq 1$.

initialize: $w_1 = w_0, \bar{g}_0 = 0$.

for $t = 1, 2, 3, \dots$ **do**

1. Given the function f_t , compute a subgradient $g_t \in \partial f_t(w_t)$.
2. Update the average subgradient \bar{g}_t :

$$\bar{g}_t = \frac{t-1}{t} \bar{g}_{t-1} + \frac{1}{t} g_t \quad (5)$$

3. Compute the next iterate w_{t+1} :

$$w_{t+1} = \arg \min_w \left\{ \langle \bar{g}_t, w \rangle + \Psi(w) + \frac{\beta_t}{t} h(w) \right\} \quad (6)$$

end for

Regularized Dual Averaging

- For simplicity, we do not need $h(w)$

Mixed ℓ_1/ℓ_2^2 -regularization. Let $\Psi(w) = \lambda\|w\|_1 + (\sigma/2)\|w\|_2^2$ with $\lambda, \sigma > 0$. Then

$$w_{t+1}^{(i)} = \begin{cases} 0 & \text{if } |\bar{g}_t^{(i)}| \leq \lambda, \\ -\frac{1}{\sigma} \left(\bar{g}_t^{(i)} - \lambda \operatorname{sign}(\bar{g}_t^{(i)}) \right) & \text{otherwise,} \end{cases} \quad i = 1, \dots, n.$$

Of course, setting $\lambda = 0$ gives the algorithm for pure ℓ_2^2 -regularization.

- Different from L1 regulariztion

$$f(w_i) = w_i - \eta \nabla_1 L(w_i, z_i) - \eta g \operatorname{sgn}(w_i),$$

- this method does not produce sparse weights online simply because only in very rare cases do two floats add up to 0

Computational Experiments

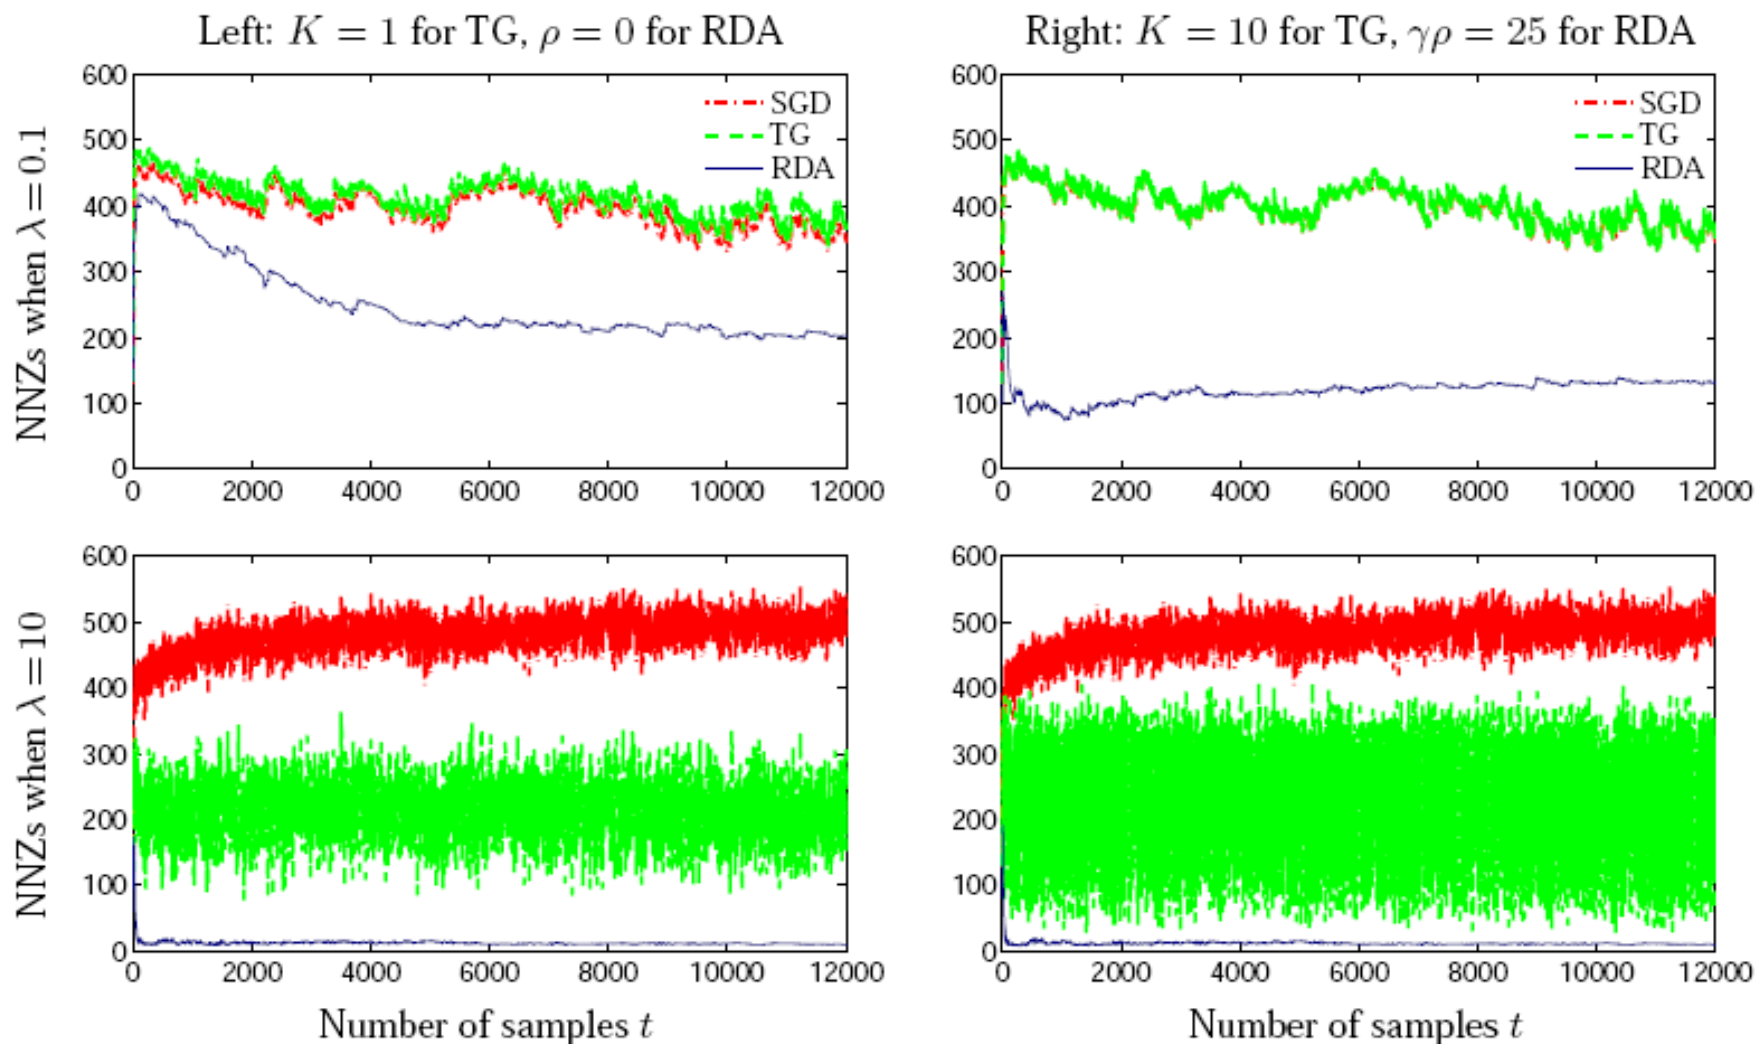


Figure 2: Number of non-zeros (NNZs) in $w(t)$ for the three online algorithms (classifying 6 and 7).

TFRL-Proximal

- General online gradient descent performs the update

$$\mathbf{w}_{t+1} = \mathbf{w}_t - \eta_t \mathbf{g}_t,$$

- Lazy representation of the model coefficients \mathbf{w} , using the "Follow The (Prox-imately) Regularized Leader" algorithm

$$\mathbf{w}_{t+1} = \arg \min_{\mathbf{w}} \left(\mathbf{g}_{1:t} \cdot \mathbf{w} + \frac{1}{2} \sum_{s=1}^t \sigma_s \|\mathbf{w} - \mathbf{w}_s\|_2^2 + \lambda_1 \|\mathbf{w}\|_1 \right)$$

- Where $\mathbf{g}_{1:t} = \sum_{s=1}^t \mathbf{g}_s$, $\sigma_{1:t} = \frac{1}{\eta_t}$.

- Different from RDA

$$\Psi(w) = \lambda \|w\|_1 + (\sigma/2) \|w\|_2^2$$

TFRL-Proximal

- New gradient update

$$\mathbf{w}_{t+1} = \begin{cases} 0 & \text{if } |z_t| \leq \lambda_1 \\ -\eta_t(z_t - \text{sgn}(z_t)\lambda_1) & \text{otherwise.} \end{cases}$$

- Where

$$z_{t-1} = \mathbf{g}_{1:t-1} - \sum_{s=1}^{t-1} \sigma_s \mathbf{w}_s$$

$$z_t = z_{t-1} + g_t + \left(\frac{1}{\eta_t} - \frac{1}{\eta_{t-1}}\right)\omega_t$$

- Per-Coordinate Learning Rates

$$\eta_{t,i} = \frac{\alpha}{\beta + \sqrt{\sum_{s=1}^t g_{s,i}^2}},$$

TFRL-Proximal

Algorithm 1 Per-Coordinate FTRL-Proximal with L_1 and L_2 Regularization for Logistic Regression

Input: parameters $\alpha, \beta, \lambda_1, \lambda_2$

$(\forall i \in \{1, \dots, d\})$, initialize $z_i = 0$ and $n_i = 0$

for $t = 1$ **to** T **do**

 Receive feature vector \mathbf{x}_t and let $I = \{i \mid x_i \neq 0\}$

 For $i \in I$ compute

$$w_{t,i} = \begin{cases} 0 & \text{if } |z_i| \leq \lambda_1 \\ -\left(\frac{\beta + \sqrt{n_i}}{\alpha} + \lambda_2\right)^{-1} (z_i - \text{sgn}(z_i)\lambda_1) & \text{otherwise.} \end{cases}$$

 Predict $p_t = \sigma(\mathbf{x}_t \cdot \mathbf{w})$ using the $w_{t,i}$ computed above

 Observe label $y_t \in \{0, 1\}$

for all $i \in I$ **do**

$g_i = (p_t - y_t)x_i$ *#gradient of loss w.r.t. w_i*

$z_i \leftarrow z_i + g_i - \frac{w_{t,i}}{\alpha} \left(\sqrt{n_i + g_i^2} - \sqrt{n_i} \right)$

$n_i \leftarrow n_i + g_i^2$

end for

end for

Empirical Results

	Num. Non-Zero's	AucLoss Detriment
FTRL-PROXIMAL	baseline	baseline
RDA	+3%	0.6%
FOBOS	+38%	0.0%
OGD-COUNT	+216%	0.0%

大纲

- 通用在线学习算法
 - 在线学习算法的应用背景
 - SGD、mini-batch算法简介 (Andrew NG)
 - SGD算法的收敛特性 (T. Zhang)
- SGD稀疏化方法
 - 稀疏化方法概要
 - Truncate Gradient
 - Regularized Dual Averaging
 - TFRL-Proximal
- 大规模在线学习系统实践
 - 节省内存
 - 有效地模型评估
 - 其它问题的探讨

Saving Memory

□ Subsampling Training Data

- Any query for which at least one of the ads was clicked.
- A fraction $r \in (0, 1]$ of the queries where none of the ads were clicked.


□ Bias calibrating

$$\omega_t = \begin{cases} 1 & \text{event } t \text{ is in a clicked query} \\ \frac{1}{r} & \text{event } t \text{ is in a query with no clicks.} \end{cases}$$

□ Applicable for positive and negative samples non-balance

Saving Memory

- Probabilistic Feature Inclusion
 - Poisson Inclusion
 - Bloom Filter Inclusion

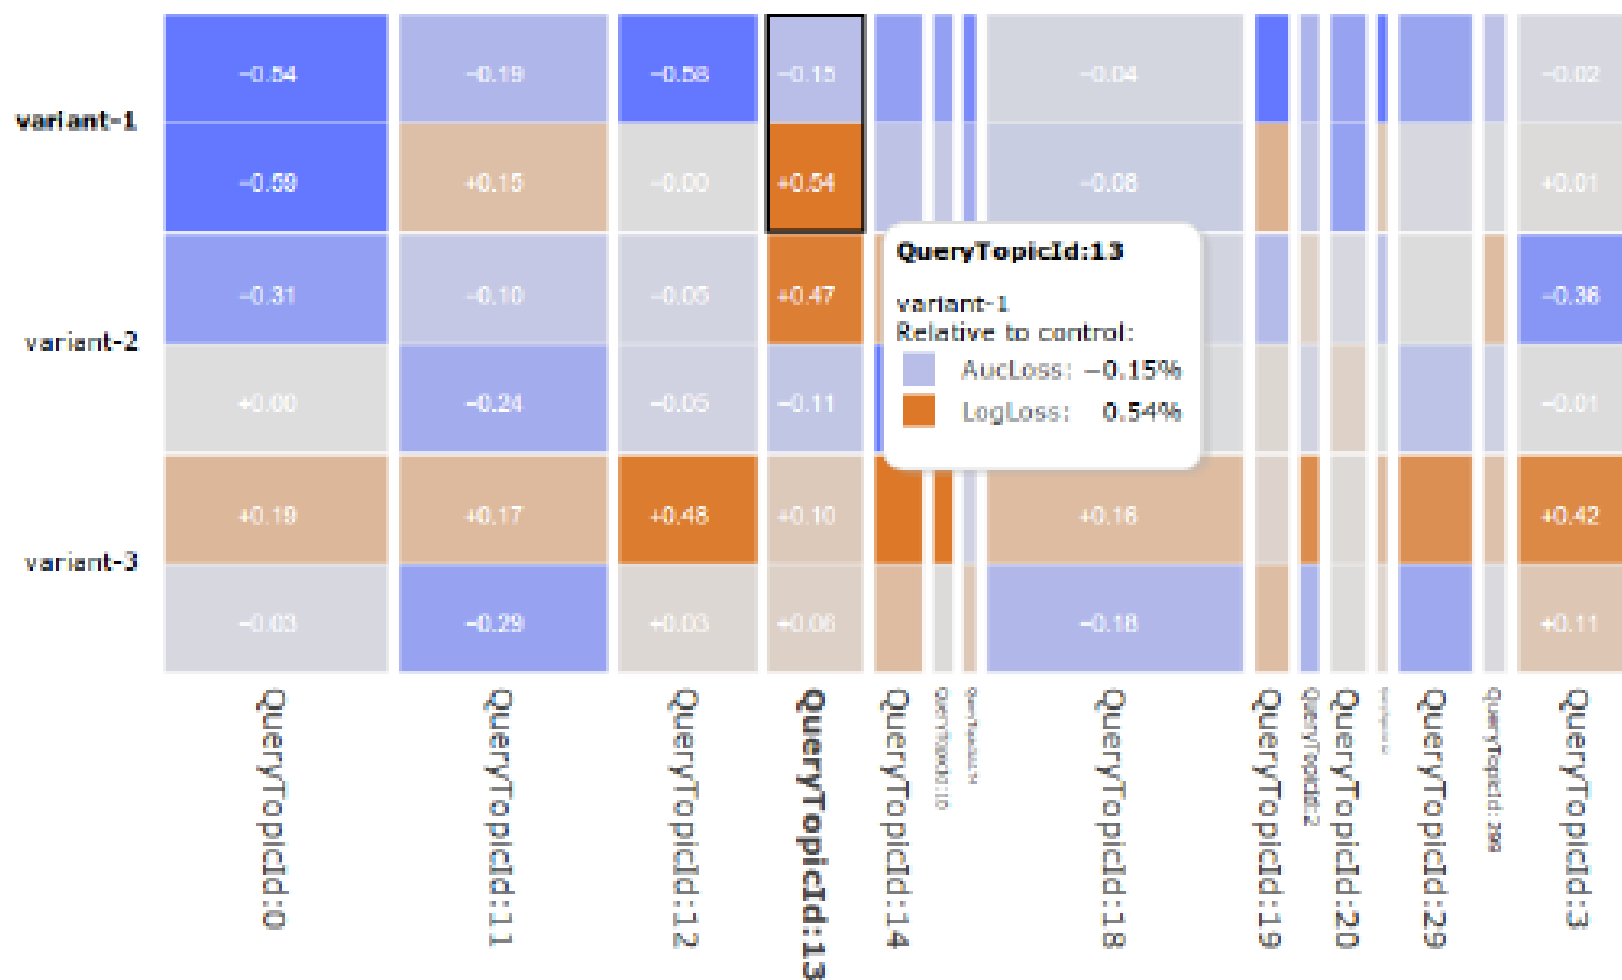


METHOD	RAM Saved	AucLoss Detriment
BLOOM ($n = 2$)	66%	0.008%
BLOOM ($n = 1$)	55%	0.003%
POISSON ($p = 0.03$)	60%	0.020%
POISSON ($p = 0.1$)	40%	0.006%

- Grafting maybe better

- Scaling up machine learning: Parallel and distributed approaches.2011. P361

Evaluating Model Performance



其它问题的探讨

- 计算平台的选择
 - Hadoop Map/Reduce 效率低
 - MPI 控制Streaming Data比较困难，节点之间耦合性高
 - 流式计算（Storm、S4）存在丢数据的风险，计算恢复困难
- 系统稳定性如何控制？

参考文献

- M. Zinkevich. Online convex programming and generalized infinitesimal gradient ascent. In Proceedings of ICML-03, pages 928–936, 2003.
- T. Zhang. Solving large scale linear prediction problems using stochastic gradient descent algorithms. In Proceedings of ICML-04, pages 919–926, 2004.
- J. Langford, L. Li, and T. Zhang. Sparse online learning via truncated gradient. JMLR, 10, 2009.
- L. Xiao. Dual averaging method for regularized stochastic learning and online optimization. In NIPS, 2009.
- Ad Click Prediction: a View from the Trenches. H. Brendan McMahan, et al, ICML 2013
- R. Bekkerman, M. Bilenko, and J. Langford. Scaling up machine learning: Parallel and distributed approaches. 2011.